

Classroom 50: Every Graded Item

One page with the accept link, the marks and the deadline for everything that is graded by tests this semester, and the six commands you need to work an item from accept to report.

Prepared by **Dr. Abhijith Anandakrishnan**

Assistant Professor, Amrita School of AI

SECTIONS

AIM A (aim-a-hpc) and AIM B (aim-b-hpc)

ORGANISATION

Amrita-School-of-AI

TOOL

gh student

UPDATED

29 Sep 2026

OFFERING

Odd Semester 2026-27

Read this first

Every lab, coding exercise and assignment in this course is a **Classroom 50** item: a private GitHub repository created for you the moment you accept it, with an autograder that reports on every submission within a couple of minutes. There is nothing to email and nothing to upload. You accept, you work, you submit, you read the report, you submit again if you want to.

Use the links for **your own section**. Accepting an item in the other section creates a repository nobody will grade.

PITFALL · YOUR GITHUB USERNAME MUST BE THE ONE ON THE ROSTER

You accepted an invitation to the `Amrita-School-of-AI` organisation at your college address. Only that account can accept items. If `gh student whoami` prints a different name, log out and log in with the right account before accepting anything.

Part 1: One-time setup

Do this once, on the machine you will work from (your laptop or the cluster login node). It takes about fifteen minutes if your SSH key already works; see `student-guide.pdf` Part 2 if it does not.

```
# 1. GitHub CLI (Debian/Ubuntu; see cli.github.com for other systems)
sudo apt install gh

# 2. Log in, then add the student tool and log in with the scopes it needs
gh auth login                # choose GitHub.com, SSH, browser login
gh extension install foundation50/gh-student
gh student login
gh student whoami            # must print your roster username

# 3. Compiler, OpenMP and MPI for local checks
sudo apt install build-essential openmpi-bin libopenmpi-dev libopenblas-dev
```

On the cluster, `gh` and the compilers are already installed; run only steps 2.

Part 2: Working an item, start to finish

Six commands. The example is Exercise 06 in section A; substitute your section and the item you are working on.

```
# 1. Accept: creates Amrita-School-of-AI/aim-a-hpc-ex06-openmp-reduction-<you>
gh student accept Amrita-School-of-AI aim-a-hpc ex06-openmp-reduction
#   or open the accept link from the tables below in a browser and click Accept

# 2. Clone it (the https:// form asks for a password and fails; use one of these)
gh repo clone Amrita-School-of-AI/aim-a-hpc-ex06-openmp-reduction-<you>
cd aim-a-hpc-ex06-openmp-reduction-<you>

# 3. Read README.md. It is the whole statement, including the marks table.

# 4. Work, and run the public tests as often as you like
make check

# 5. Commit and submit
git add -A
git commit -m "say what you changed"
gh student submit

# 6. Read the report, fix, repeat. Before the second submit:
git pull --rebase origin main
```

Where the report is

About a minute after `gh student submit`:

- the commit on GitHub gets a **green tick or a red cross**;
- a **Release** appears on your repository, named `submit/<time>-<sha>`. Its text is the grader's report: one line per check, PASS or FAIL, the marks for it, and the reason for every failure;
- the same text is under the run's **Summary** in the Actions tab;
- on classroom50.org, **My submission** shows the score and links to the release.

If there is no release after two minutes, open the Actions tab. A red run whose log ends in a grader error, rather than a failing test, is a problem on my side: send me the link.

What you may and may not edit

The README of each item names the files that are yours: usually `src/*.c`, sometimes `answers.yaml` or a CSV under `results/`. Everything else is fixed. `.gitignore` and `.github/` are restored on every submit; `tests/` holds copies of the public tests and the grader uses its own copy, so changing them changes nothing. The reference functions marked *given* in a header must stay as they are: the tests compare your versions against them.

How marks are earned

KIND OF CHECK	WHAT PASSES IT
Compiles	<code>-Wall -Wextra -Werror</code> . Fails, and nothing after it runs.
Required constructs	the mechanism the item teaches is in the source; comments do not count
Public tests	the tests <code>make check</code> ran
Hidden tests	the same behaviour on inputs you have not seen: odd sizes, tiny sizes, awkward rank counts
Determinism	identical output over repeated runs at several thread counts
Speedup	the faster version is faster by a modest factor on a two-core machine; only checked once the code is correct
Analysis	your program computes speedup, efficiency, Karp-Flatt or roofline quantities exactly on hidden data
Committed results	the CSV you produced on the cluster has the right columns, rows and ranges
Answers	numbers in <code>answers.yaml</code> against a key, with a tolerance

Some marks cannot be earned on the grading machine. Items marked *GPU cluster* in the tables have those checks run on `asaicompute` after the deadline and added to your total; the report tells you how many marks are pending. The four assignments carry 4 marks for a short written section, `EXPLANATION.md`, graded against the rubric printed in the README. Everything else is marked by the tests, and the last submission before the deadline is the one that counts.

Part 3: The items

Accept links are per section. **A** is `aim-a-hpc`, **B** is `aim-b-hpc`. Due times are 23:59 IST on the day shown. Late submissions are marked out of 80 percent within 48 hours and not accepted after that, as in the handbook.

Every item is open now. The deadlines are spread evenly from mid-September to 23 October, the last working day of the semester, in syllabus order, so that one item falls due roughly every other day. Nothing stops you from working ahead: accept whatever you like now, and submit early. The order in the tables is the order the material is taught, and the sensible order to do them in.

Everything that needs the GPU cluster now closes on **Tuesday 20 October**: Labs 2, 3 and 8, Exercises 17 and 19, and Assignments 1, 2 and 3, alongside Exercise 20 which was already on that date. You have no accounts on `asaicompute` yet, and it would be unfair to hold you to a deadline for work you cannot start. Details of how to log in will follow separately; the rest of each of those items can be done now.

Lab 1 and Exercises 01 and 02 close on **Monday 28 September** as already announced.

PITFALL · ONE ITEM EVERY TWO DAYS IS THE PACE, NOT A SUGGESTION

The dates are close together on purpose. An item left for the weekend before the deadline collides with the two behind it. Do each one the week its material is covered and the schedule is light; leave them and it is not.

Warm-up (no marks)

ITEM	NAME	MARKS	ACCEPT
<code>a0-git-warmup</code>	A0: Git and GitHub warm-up	0	A · B

Do A0 first. It is the whole loop above on a task that cannot hurt you, and it tells us both that your account, your tools and the grader agree.

Coding exercises (2% of the course, pooled)

ITEM	NAME	MARKS	DUE (IST)	ACCEPT	GRADED WHERE
ex01-performance-metrics	Ex 01: Performance metrics and the roofline	10	Mon 28 Sep, 23:59	A · B	runner only
ex02-cache-blocking	Ex 02: Cache blocking	12	Mon 28 Sep, 23:59	A · B	runner only
ex06-openmp-reduction	Ex 06: Parallel reduction with OpenMP	10	Mon 28 Sep, 23:59	A · B	runner only
ex07-openmp-load-balance	Ex 07: Load imbalance and the schedule clause	12	Tue 29 Sep, 23:59	A · B	runner only
ex08-openmp-histogram	Ex 08: Parallel histogram	12	Thu 01 Oct, 23:59	A · B	runner only
ex11-mpi-halo-stencil	Ex 11: MPI halo exchange	10	Thu 08 Oct, 23:59	A · B	runner only
ex12-mpi-collectives	Ex 12: MPI collectives	11	Sat 10 Oct, 23:59	A · B	runner only
ex17-cuda-vector-add	Ex 17: First CUDA kernels	10	Tue 20 Oct, 23:59	A · B	7 on the GPU cluster
ex19-cuda-tiled-matmul	Ex 19: Shared-memory tiling	12	Tue 20 Oct, 23:59	A · B	7 on the GPU cluster
ex20-cuda-reduction	Ex 20: Parallel reduction on a GPU	10	Tue 20 Oct, 23:59	A · B	7 on the GPU cluster

Each exercise is one file in `src/`. The starters compile and are correct serial code; the marks are for the parallel or optimised version the statement asks for, and for keeping it correct on the hidden inputs.

Lab sheets (2% of the course, pooled)

ITEM	NAME	MARKS	DUE (IST)	ACCEPT	GRADED WHERE
lab01-timing-and-memory	Lab 1: Timing and memory	10	Mon 28 Sep, 23:59	A · B	runner only
lab02-slurm-and-the-cluster	Lab 2: SLURM and the cluster	10	Tue 20 Oct, 23:59	A · B	runner only
lab03-scaling-study-and-roofline	Lab 3: Scaling study and roofline	10	Tue 20 Oct, 23:59	A · B	runner only
lab04-openmp-and-races	Lab 4: OpenMP and races	10	Sat 26 Sep, 23:59	A · B	runner only
lab05-optimising-matrix-multiply	Lab 5: Optimising matrix multiply	10	Sat 03 Oct, 23:59	A · B	runner only
lab06-mpi-basics-and-deadlock	Lab 6: MPI basics and deadlock	10	Tue 06 Oct, 23:59	A · B	runner only
lab07-domain-decomposition	Lab 7: Domain decomposition	10	Sun 11 Oct, 23:59	A · B	runner only
lab08-first-cuda-kernels	Lab 8: First CUDA kernels	10	Tue 20 Oct, 23:59	A · B	3 on the GPU cluster
lab09-cuda-profiling-and-roofline	Lab 9: CUDA profiling and roofline	10	Wed 21 Oct, 23:59	A · B	runner only

The lab sheets keep their original teaching text; the last section of each README, *What is graded*, says which deliverables the tests look at. Measurements that need the cluster are committed as CSVs and shape-checked; the numbers themselves are yours to interpret.

Assignments (6% of the course, pooled)

ITEM	NAME	MARKS	DUE (IST)	ACCEPT	GRADED WHERE
<code>a1-measure-before-you-optimize</code>	Assignment 1: Measure before you optimize	30	Tue 20 Oct, 23:59	A · B	4 by rubric
<code>a2-openmp-scaling</code>	Assignment 2: OpenMP scaling	30	Tue 20 Oct, 23:59	A · B	4 by rubric
<code>a3-mpi-decomposition</code>	Assignment 3: MPI decomposition	30	Tue 20 Oct, 23:59	A · B	4 by rubric
<code>a4-gpu-and-ai</code>	Assignment 4: GPU and AI	30	Fri 23 Oct, 23:59	A · B	14 on the GPU cluster; 4 by rubric

Thirty marks each: 26 from tests, 4 from `EXPLANATION.md` (at most 250 words, rubric in the README). The kernel is fixed for everyone so that the tests are the same for everyone: Jacobi stencil, CSR sparse matrix-vector product, 2-D heat diffusion with MPI, and a CUDA matrix transpose.

Part 4: When something goes wrong

SYMPTOM	WHAT IT MEANS
<code>gh student accept</code> says the assignment is not found	You are in the other section, or the item is not released yet.
<code>git clone https://...</code> asks for a password	Use <code>gh repo clone</code> or the <code>git@github.com:</code> address.
<code>make check</code> passes, the report shows hidden tests failing	Your code is right for the sizes you tried. Read the failing test's name: it names the case.
<code>! [rejected] main -> main (fetch first)</code> on push	Submitting added a commit. <code>git pull --rebase origin main</code> , then push again.
Speedup check says <i>skipped: ... failed</i>	Speed is only measured once correctness passes. Fix the tests first.
The report says marks are graded on the cluster	Normal for CUDA items. They arrive after the deadline.
No release two minutes after submitting	Open the Actions tab. If the run itself is red, send me the link.
<code>nvcc: command not found</code> on your laptop	CUDA items build on the cluster (<code>srun --gres=gpu:1</code>), not on a laptop without a GPU.